Question 4.   [6 MARKS]   Question 1

A program issues the following system calls on the original Fast File System (FFS) file system. The file system block size is 4KB and the size of an inode is 64 bytes. Assume that the file system is already mounted when the program starts and that all system calls succeed.

```
char buf[4096];
int fd = open("/a/b/c", 0); // open in read-only mode
lseek(fd, 1034*4096, 0);     // seek to position (1034*4096) from start of file
read(fd, buf, 4096);         // read 4k of data from file
```

Assume that the file buffer cache is empty. State the *minimum* number of the following types of blocks that will be read when the program above is run. Explain your answer for each block type.

1. Inode block(s)

2. Directory block(s)

3. Indirect block(s) (include single, double or triple indirect)

4. Other data block(s)

Question 5. File system consistency [9 MARKS] **Question 2**

On an ext2 or FFS file system, consider the operation of creating a new directory in an existing directory. Assume the existing directory occupies one block and there is enough space to add a new entry. Assume the existing directory inode and the new directory inode are in different disk blocks.

Part (a) [2 MARKS]

Which of the following blocks must be updated? Check all that apply.

☐ inode bitmap
☐ new directory inode
☐ new directory inode
☐ new directory data block
☐ existing directory inode
☐ existing directory data block

Part (b) [2 MARKS]

What data is updated in the existing directory inode?

last modified time, number of links

Part (c) [5 MARKS]

In each of the remaining questions, check all of the boxes that most closely explain what happens if a crash occurs after updating only the block(s) specified.

Inode Bitmap and Data Block Bitmap

☐ No inconsistency (it simply appears that the operation was not performed)
☐ Data leak (data block is lost for any future use)
☐ Inode leak (Inode is lost for any future use)
☐ Inconsistent inode data (Some inode field does not match what is stored in data blocks)
☐ Multiple file paths may point to same inode
☐ Something points to garbage

New Directory Inode

☐ No inconsistency (it simply appears that the operation was not performed)
☐ Data leak (data block is lost for any future use)
☐ Inode leak (Inode is lost for any future use)
☐ Multiple file paths may point to same inode
☐ Inconsistent inode data (Some inode field does not match what is stored in data blocks)

☐ Something points to garbage

Inode Bitmap, Data Block Bitmap, Existing Directory data, New Directory inode, and New Directory data

☐ No inconsistency (it simply appears that the operation was not performed)

☐ Data leak (data block is lost for any future use)

☐ Inode leak (Inode is lost for any future use)

☐ Multiple file paths may point to same inode

☐ Inconsistent inode data (Some inode field does not match what is stored in data blocks)

☐ Something points to garbage

Inode Bitmap and New Directory inode

☐ No inconsistency (it simply appears that the operation was not performed)

☐ Data leak (data block is lost for any future use)

☐ Inode leak (Inode is lost for any future use)

☐ Multiple file paths may point to same inode

☐ Inconsistent inode data (Some inode field does not match what is stored in data blocks)

☐ Something points to garbage

New Directory inode, Existing Directory inode, and Existing Directory data

☐ No inconsistency (it simply appears that the operation was not performed)

☐ Data leak (data block is lost for any future use)

☐ Inode leak (Inode is lost for any future use)

☐ Multiple file paths may point to same inode

☐ Inconsistent inode data (Some inode field does not match what is stored in data blocks)

☐ Something points to garbage

Part (e) [1 MARK]   Question 3

Indexed-based file systems suffer from external fragmentation because blocks of the file may be scattered across the disk.

Part (f) [1 MARK]

Extent-based file systems may require more disk block accesses than indexed-based files systems for random access because it may need to traverse an extent to get to a particular byte in the file.

Q7. [16 marks] File systems [25 minutes] Question 4

Consider the following ext2 file system, as discussed in class (you can ignore the absence of the superblock and block group regions).

Inode bitmap (IB)	Data bitmap (DB)	Inode table (I)	Data region (D)

1) [5 marks] A user creates a new (empty) file A in a given directory. You can assume that the file name is not taken, that the directory does exist, and that the file system is consistent at this point.

  i) Which of the regions above must be updated for the operation to be complete?

  ii) Imagine that a crash can occur during this operation, and that we have no tool for checking file system consistency. In which order should the updates happen, such that we minimize the potential negative impact on the file system? Explain your rationale **in detail**.

2) [5 marks] Imagine a deletion operation for a file B. You can assume that the file exists and that no file system inconsistency is present at this time.

  i) Which of the regions above must be updated for the operation to be complete?

  ii) Imagine that a crash can occur during this operation, and that we have no tool for checking file system consistency. In which order should the updates happen, such that we minimize the potential negative impact on the file system? Explain your rationale **in detail**.

3) [3 marks] What was the main motivation behind the design of the Log-structured File System (LFS)? How are updates carried out on LFS? In your explanation, compare to how updates are carried out in the Fast File System (FFS).

4) [3 marks] Describe what is the challenge in locating data and metadata on disk in LFS. Additionally, explain in detail how this is done in practice.

Q7. [16 marks] File systems [25 minutes] Question 5

1) [4 marks] Consider that you're implementing an ext2_restore program similar to the one from A4. You are attempting to recover the file /home/bart/file1 and go through the "gaps" in the directory entries of the "bart" directory. You end up finding a removed directory entry with a valid (non-zero) inode and the name "file1", then you go to the inode bitmap and find that the inode is in use. What can you infer from this? Explain **in detail**, **everything** that you can infer from this.

2) [5 marks] The fsck tool typically runs at boot time while the file system is not in active use by any user. User Homer finds that fsck is very slow and does not want to wait for it to run at boot time. Instead, after boot time, he runs fsck manually in a terminal, using a set of flags which enable fsck to run "online" (during normal operation rather than at boot time).
Can any problem occur in this situation? If a problem could occur, explain *in detail* what could happen with this approach. If not, then explain *in detail* why the approach is fine.

3) [5 marks] Metadata journaling provides reasonable but not full consistency guarantees like full (data and metadata) journaling. Explain in detail why this is the case, and describe a scenario where metadata journaling fails to provide full consistency.
Hint: Write down the steps of metadata journaling and compare to the steps of full journaling.

4) [2 marks] Describe a method that SSDs use to prevent wearout of the flash cells.

Question 8. File Systems [15 MARKS]   Question 6

A defragmentation program reorganizes the blocks on a disk to maximize the contiguous free space on the disk. You can think of the blocks on the disk as an array, where the block number is an index into the array. Consider the following two algorithms:

- Alg. A: Free blocks are filled by used blocks from the end of the disk.
- Alg. B: Blocks are moved around until all the files in the block group are stored contiguously.

Part (a) [1 MARK]

Which of the two algorithms would complete faster on average? Explain your thinking.

Part (b) [2 MARKS]

Which of the two algorithms would lead to better performance in terms of time to access blocks during normal operation? Assume we are using a magnetic disk. Explain clearly.

Part (c) [2 MARKS]

Consider an inode based-file system like ext2. What changes, if any, are made to the inodes? Explain your answer.

Part (d) [2 MARKS]

Consider an inode based-file system like ext2. What changes, if any, would need to be made to the directory metadata by the defragmentation algorithm? Explain your answer.

Part (e) [2 MARKS]

Consider a linked file system structure like FAT. What changes, if any, are made to the FAT (File Allocation Table) by the defragmentation algorithm? Explain your answer.

Part (f) [2 MARKS]

Consider a linked file system structure like FAT. What changes, if any, are made to the directory metadata by the defragmentation algorithm? Explain your answer.

Part (g) [4 MARKS]

Solid-state hard drives have uniform access time to all parts of the drive.

i– [2 MARKS] Explain one reason why you might not want to move blocks around so that they are close to other blocks from the same file.

ii– [2 MARKS] Explain one reason why you might still want blocks from the same file to be close together on the drive.

Question 10. File system consistency [7 MARKS]   Question 7

On an ext2 file system, consider the operation of creating a new empty file in an existing directory. Assume the directory occupies one block and there is enough space to add a new entry. Assume the directory inode and the file inode are in different disk blocks.

Which of the following blocks must be updated? Check all that apply.

- ☐ inode bitmap
- ☐ data bitmap
- ☐ file inode
- ☐ directory inode
- ☐ directory data block

In each of the remaining questions, check the box that most closely explains what happens if a crash occurs after updating only the block(s) specified.

Bitmap

- ☐ No inconsistency (it simply appears that the operation was not performed)
- ☐ Data leak (data block is lost for any future use)
- ☐ Inode leak (Inode is lost for any future use)
- ☐ Multiple file paths may point to same inode
- ☐ Something points to garbage
- ☐ Multiple of the problems listed above

File inode

- ☐ No inconsistency (it simply appears that the operation was not performed)
- ☐ Data leak (data block is lost for any future use)
- ☐ Inode leak (Inode is lost for any future use)
- ☐ Multiple file paths may point to same inode
- ☐ Something points to garbage
- ☐ Multiple of the problems listed above

Directory inode and directory data

- ☐ No inconsistency (it simply appears that the operation was not performed)
- ☐ Data leak (data block is lost for any future use)
- ☐ Inode leak (Inode is lost for any future use)
- ☐ Multiple file paths may point to same inode
- ☐ Something points to garbage
- ☐ Multiple of the problems listed above

Bitmap and file inode

☐ No inconsistency (it simply appears that the operation was not performed)

☐ Data leak (data block is lost for any future use)

☐ Inode leak (Inode is lost for any future use)

☐ Multiple file paths may point to same inode

☐ Something points to garbage

☐ Multiple of the problems listed above

Bitmap and directory inode and directory data

☐ No inconsistency (it simply appears that the operation was not performed)

☐ Data leak (data block is lost for any future use)

☐ Inode leak (Inode is lost for any future use)

☐ Multiple file paths may point to same inode

☐ Something points to garbage

☐ Multiple of the problems listed above

File inode and directory inode and directory data

☐ No inconsistency (it simply appears that the operation was not performed)

☐ Data leak (data block is lost for any future use)

☐ Inode leak (Inode is lost for any future use)

☐ Multiple file paths may point to same inode

☐ Something points to garbage

☐ Multiple of the problems listed above

Question 8

Q7. [16 marks] File systems [25 minutes]

1) [6 marks] Assume that your course grades are stored in your file system in a regular file called "`mygrades.txt`" and that you are operating in the directory where this file is located.

Consider the following two *independent* scenarios:

S1) You run the following commands:
```
ln mygrades.txt ../backups/mygrades-copy.txt
cd ../backups
ln -s mygrades-copy.txt mygrades-copy2.txt
```

S2) You run the following commands:
```
ln -s mygrades.txt ../backups/mygrades-copy.txt
cd ../backups
ln mygrades-copy.txt mygrades-copy2.txt
```

Answer the following questions below, *for each of the scenarios above*:
a) What happens if you remove the file `mygrades-copy.txt`? Explain in detail the implications.
b) What happens if you remove the file `mygrades.txt`? Explain in detail the implications.
Note: The two questions are *independent*. That is, for b) do not assume that `mygrades-copy.txt` has been removed.

2) [3 marks] In ext2, what happens when we remove a file which happens to be stored as the first directory entry in the second data block of its parent directory.

3) [3 marks] How does fsck ensure that the data bitmaps match the data blocks which are in use by files and directories? Explain how this invariant is checked and what repair actions are taken.

4) [4 marks] You are consulting for a storage company that needs to ensure specific goals for their data storage system. Their **first goal** is redundancy in case of complete disk failures. At the same time, an **equally important goal** is good data bandwidth. The company currently stores their data in a RAID0 (striped) and you are asked to provide suggestions that would better suit the company's goals in terms of data storage. What would you recommend?

Question 6. File Systems [12 MARKS] Question 9

Suppose you have a Unix file system where the block size is 2K bytes, a disk address is 32 bits, and an i-node contains the disk addresses of:

- the first 12 blocks of file
- a single indirect block, and
- a double indirect block

Part (a) [4 MARKS] What is the largest file that can be represented by an i-node?

Part (b) [4 MARKS] Suppose that there is a file in the system that is of the maximum size (from your answer above). Assuming the i-node for this file is not in memory yet, how many disk accesses in total are required for a program to read the entire file in sequential order? Assume that one entire block is read into memory at a time and reading data from a block that is already in memory does not require a disk access. Account for every disk access related to reading this file (but dont count any paging in the virtual memory system), and show all your work.

Part (c) [4 MARKS] briefly explain 2 main ways of implementing directoreies

Question 6.   File Systems [12 MARKS]   **Question 10**

Part (a)   [4 MARKS] Suppose that on a different computer, the OS uses i-nodes and each disk block is 4KB (2^{12} bytes). Assume that an i-node contains 12 direct block numbers (disk addresses) and the block numbers for one indirect block, one double indirect block, and one triple indirect block. Assume also that a block number is 4 bytes. What is the largest possible file on that computer (assuming the disk is large enough)?

Part (b)   [4 MARKS] On the computer from part above, suppose you wanted to create a new file of the maximum size (that you computed for the answer to (a)). How many free blocks on the disk would there need to be in order to create that file?

Part (c)   [4 MARKS] What's the difference between hardlink and softlink?