

CSC 369 Worksheet 5 Solution

August 18, 2020

1. I need to run randomly-generated problems with two jobs and two queues using file `mlfq.py` with I/O turned off, and compute the MLFQ execution trace for each.

Using the command `./mlfq.py -s 1 -m 10 -n 2 -j 2 -M 0`, we have

```
Job List:
Job  0: startTime  0 - runTime  2 - ioFreq  0
Job  1: startTime  0 - runTime  7 - ioFreq  0
```

with

- allotment time for queue 1 is 1
- quantum length for queue 1 is 10
- allotment time for queue 0 is 1
- quantum length for queue 0 is 10
- no priority boost

,

the execution trace is:

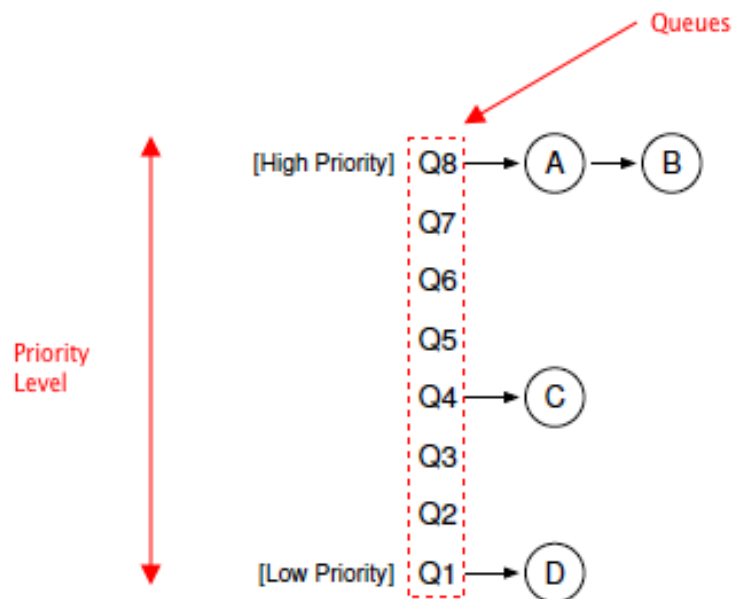
Notes

- `-m 10` sets the maximum runtime of a job to 10
- `-M 0` turns off I/O in `mlfq.py`
- `-n 2` sets number of queues to 2
- `-j 2` sets number of jobs to 2
- **Multi-level Feedback Queue (MLFQ):**

- Is one of the most well-known approaches to scheduling
- Does two things:
 - a) Optimizes turnaround time
 - b) Minimizes response time
- Uses **priority level** and **Queues** to achieve it's goal

- **MLFQ Basic Rules:**

- Jobs on same queue → Same priority
- **Rule 1:** If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't)
- **Rule 2:** If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR

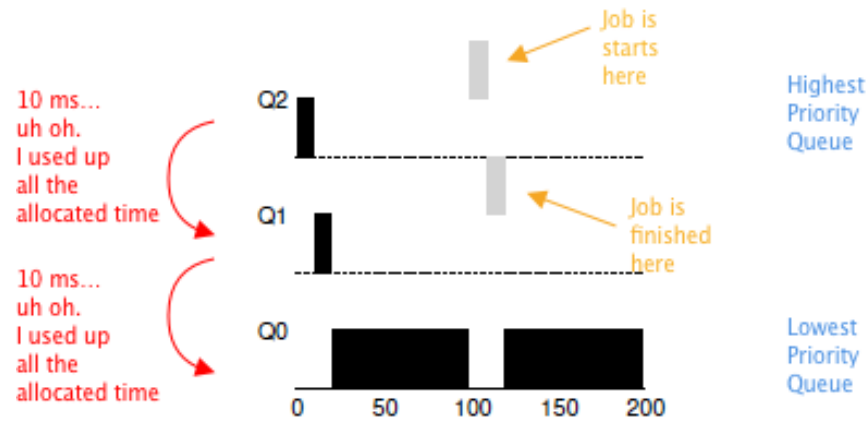


- **Attemp #1: How to Change Priority**

- **Rule 3:** When a job enters the system, it is placed at the highest priority (the topmost queue)
- **Rule 4a:** If a job uses up an entire time slice while running, its' priority is reduced (i.e. it moves down on queue).
- **Rule 4b:** If a job gives up the CPU before the time slice is up, it stays at the same priority level (e.g I/O Operation)
 - * Means that the shifting down of priority level only depends on CPU time

Example (Along Came a Short Job):

- 1) A job A enters system
- 2) Job is placed on highest Queue Q_2
- 3) After time-slice (e.g. 10 ms) in Q_2 , A is placed on lower queue Q_1
- 4) After time-slice in Q_1 , A is placed in lowest priority queue Q_0



- **Attemp #2: The Priority Boost**

- **Rule 5:** After some time period S , move all the jobs in the system to the topmost queue.
- * This is to prevent starvation (i.e. a job never being run)

- **Attempt #3: Better Accounting (Fix of Attempt # 1)**

- Is to prevent programmers from gaming (i.e. tricking) the CPU so all programs get a fair share of allotment time
- **Rule 4:** Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced (it moves down one queue).