

# 1 User Mode

- Is restricted
- Executing code has no ability to *directly* access hardware or reference memory <sup>[1]</sup>
- Crashes are always recoverable <sup>[1]</sup>
- Is where most of the code on our computer / applications are executed <sup>[3]</sup>

# 2 Kernel Mode

- Is privileged (non-restricted)
- Executing code has complete and unrestricted access to the underlying hardware <sup>[3]</sup>
- Is generally reserved for the lowest-level, most trusted functions of the operating system <sup>[1]</sup>
- Is fatal to crash; it will halt the entire PC (i.e the blue screen of death) <sup>[3]</sup>

# 3 Interrupt

- Are signals sent to the CPU by external devices, normally I/O devices. <sup>[2]</sup>
- Tells the CPU to stop its current activities and execute the appropriate part of the operating system (**Interrupt Handler**). <sup>[2]</sup>
- Has three different types <sup>[2]</sup>

## 1) Hardware Interrupts

- Are generated by hardware devices to signal that they need some attention from the OS.
- May be due to receiving some data

### Examples

- \* Keystrokes on the keyboard
- \* Receiving data on the ethernet card

- May be due to completing a task which the operating system previously requested

### Examples

Transferring data between the hard drive and memory

## 2) Software Interrupts

- Are generated by programs when a system call is requested

## 3) Traps

- Are generated by the CPU itself
- Indicate that some error or condition occurred for which assistance from the operating system is needed

# 4 Content Switch

- Is switching from running a user level process to the OS kernel and often to other user processes before the current process is resumed
- Happens during a timer interrupt or system call
- Saves the following states for a process during a context switch
  - Stack Pointer
  - Program Counter
  - User Registers
  - Kernel State
- May hinder performance


# 5 System Call

### Example

- `yield()`
  - Is a system call
  - Causes the calling thread to relinquish the CPU
  - Places the current thread at the end of the run queue
  - Schedules another thread to run

## 6 Thread

- Is a lightweight process that can be managed independently by a scheduler <sup>[4]</sup>
- Improves the application performance using parallelism. (e.g. peach)



- A thread is bound to a single process
- A process can have multiple threads
- Has two types
  - **User-level Threads:**
    - \* Are implemented by users and kernel is not aware of the existence of these threads
    - \* Are represented by a program counter(PC), stack, registers and a small process control block

- \* Are small and much faster than kernel level threads
- **Kernel-level Threads:**
  - \* Are handled by the operating system directly
  - \* Thread management is done by the kernel
  - \* Are slower than user-level threads

## 7 Process

- Is a program in execution
- Is named by it's process ID or PID
- Can be described by the following states at any point in time
  - Address Space
  - CPU Registers
  - Program Counter
  - Stack Pointer
  - I/O Information

(wait. this is PCB)

- Exists in one of many different **process states**, including
  1. Running
  2. Ready to Run
  3. Blocked
- Different events (Getting Scheduled, descheduled, or waiting for I/O) transitions one of these states to the other

## 8 Signals

- Provides a way to communicate with the process
- Can cause job to stop, continue, or terminate
- Can be delivered to an application
  - Stops the application from whatever its doing
  - Runs Signal handler (some code in application to handle the signal)
  - When finished, the process resumes previous behavior

## 9 Spinlock

- Is the simplest lock to build
- Uses a lock variable
  - 0 - (available/unlock/free)
  - 1 - (acquired/locked/held)

- Has two operations

1. `acquire()`

```
boolean test_and_set(boolean *lock)
{
    boolean old = *lock;
    *lock = True;
    return old;
}

boolean lock;

void acquire(boolean *lock) {
    while(test_and_set(lock));
}
```

2. `release()`

```
void release(boolean *lock) {
    *lock = false;
}
```

- Allows a single thread to enter critical section at a time
- Spins using CPU cycles until the lock becomes available.
- May spin forever

## 10 Scheduling policies

- Are algorithms for allocating CPU resources to concurrent tasks deployed on (i.e., allocated to) a processor (i.e., computing resource) or a shared pool of processors <sup>[5]</sup>
- Are sometimes called **Discipline**
- Covers the following algorithms in textbook
  - **First In First Out**
  - **Shortest Job First**
  - **Shortest Time-to-completion First**
  - **Round Robin**
    - \* Runs job for a **time slice** or **quantum**
    - \* Each job gets equal share of CPU time
    - \* Is clock-driven <sup>[6]</sup>
    - \* Is starvation-free <sup>[7]</sup>
    - \* Must have the length of a time slice (**quantum**) as multiple of timer-interrupt period

```
void release(boolean *lock) {  
    *lock = false;  
}
```

- **Multi-level Feedback Queue**

### References

- 1) Coding Horror, Understanding User and Kernel Mode, [link](#)
- 2) Kansas State University, Basics of How Operating Systems Work, [link](#)
- 3) Kansas State University, Glossary, [link](#)
- 4) Tutorials Point, User-level threads and Kernel-level threads, [link](#)
- 5) Science Direct, Scheduling Policy, [link](#)
- 6) Guru 99: What is CPU Scheduling?, [link](#)
- 7) Wikipedia: Round-robin Scheduling, [link](#)