

1. a) 1) 4 - inode blocks. 1 for the file c, and 3 for the directories /, a, b
 - 2) 3 - directory blocks - one for root /, one for a, the other for b
 - 3) 1 - single indirect block as far as we know. The file definitely has more than 12 blocks (# of data blocks pointed by direct pointers), but less than 1036 (# of data blocks pointed by direct pointers and single indirect pointers). We are reading block 1034.
 - 4) 1 - data block for file c
- b) All of the above

Notes

- **Inode**



- Is short form of **index node**
- Describes a file system object such as file or data
- Contains all information about a file/directory, including
 - * File Type,
 - * Size
 - * Number of blocks allocated to it
 - * Protection information
 - * Time information (e.g time created, time modified)
 - * Location of data blocks residing on disk

References

- 1) Wikipedia, Inode, link
 - 2) Machanick, Philip. (2016). Teaching Operating Systems: Just Enough Abstraction. 642. 10.1007/978-3-319-47680-3_10., link
- c) Size, the location of data blocks that reside on disk

Notes

- I wonder what information about blocks inode has. Is it total number of blocks both inode and data, or just data?
 - I struggled a bit on this one. I should find an easier way to remember which information inode has
- d) • **Inode Bitmap and Data Block Bitmap**
- (b) - Data Leak
 - (c) - Inode Leak
- **New Directory Inode**
- (a) - No inconsistency
- **Inode Bitmap, Data Block Bitmap, Existing Directory Data, New Directory Inode, and New Directory Data**
- (e) - Inconsistent inode data
- **Inode Bitmap, and New Directory Inode**
- (c) - Inode leak
 - (d) - Multiple file paths may point to same inode
- **New Directory Inode, Existing Directory Inode, Existing Directory Data**
- (e) - Inconsistent inode data
 - (f) - Something points to garbage

Correct Solution

- **Inode Bitmap and Data Block Bitmap**
 - (b) - Data Leak
 - (c) - Inode Leak
- **New Directory Inode**
- (a) - No inconsistency
- **Inode Bitmap, Data Block Bitmap, Existing Directory Data, New Directory Inode, and New Directory Data**
- (e) - Inconsistent inode data
- **Inode Bitmap, and New Directory Inode**
- (c) - Inode leak
- **New Directory Inode, Existing Directory Inode, Existing Directory Data**

(d) - Multiple file paths may point to same inode
 (f) - Something points to garbage

Notes

- I wonder how system call for reading file/directory works in UNIX. Does it check for bitmap?
- I wonder how system call for deleting file/directory works in UNIX
- I wonder how system call for creatubg file/directory works in UNIX
- Learned that
 - Missing Inode Bitmap - multiple file paths may point to same inode

• **File API**

- open (create/access file)
 - * Is a system call
 - * Does three things on creation
 - 1) make structure (inode) that racks all relevant information about file
 - 2) link human readable name to the file, and put that link to a directory
 - 3) increment **reference count** in inode

* **Syntax:**

```
int fd = open("foo". O_CREAT|O_WRONLY|O_TRUNC, S_IRUSR|S_IWUSR)
```

- O_CREAT - Creates file "foo" if does not exist
- O_WRONLY - Open file for writing only (default)
- O_TRUNC - Overwrites existing file **Need example/Clarification**
- Can have multiple flags
- * Returns **file descriptor** or fd for short
 - Is an integer
 - Is used to access a file
 - Is private per process
 - Can be used to read() and write() files

Example

```

#include <fcntl.h>
...
int fd;
mode_t mode = S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH;
char *filename = "/tmp/file";
...
fd = open(filename, O_WRONLY | O_CREAT | O_TRUNC, mode);
...

```

File can be read by owner

File can also be written by owner

File can also be read by group

File can also be read by others

Means

1. File is Writable AND
2. Create file if doesn't exist AND
3. Overwrite file if exists

- (read) (read file)
- * Is a system call
- * **Syntax:**

```
ssize_t read (int fd, void *buf, size_t count)
```

- fd - file descriptor (from open())
- buf - container for the read data
- count - number of bytes to read
- * Returns number of bytes read, if successful
- * Returns 0 if is at, or past the end of file

Example

```

char buf[4096];
int fd = open("/a/b/c", 0); // open in read-only mode
lseek(fd, 1034*4096, 0);    // seek to position (1034*4096) from start of file
read(fd, buf, 4096);        // read 4k of data from file

```

System Calls	Return Code	Current Offset	
fd = open("file", O_RDONLY);	3	0	
read(fd, buffer, 100);	100	100	read continues for each call
read(fd, buffer, 100);	100	200	
read(fd, buffer, 100);	100	300	
read(fd, buffer, 100);	0	300	returns 0 if at end
close(fd);	0	-	

- write (write file)

- * Is a system call
- * Writes data out of a buffer
- * **Syntax:**

```
ssize_t write (int fd, const void * buf, size_t nbytes)
```

- fd - file descriptor
- buf - A pointer to a buffer to write to file
- nbytes - number of bytes to write. If smaller than buffer, the output is truncated

Example

```
#include <unistd.h>
#include <fcntl.h>

int main(void)
{
    int filedesc = open("testfile.txt", O_WRONLY | O_APPEND);

    if (filedesc < 0) {
        return -1;
    }

    if (write(filedesc, "This will be output to testfile.txt\n", 36) != 36) {
        write(2, "There was an error writing to testfile.txt\n", 43);
        return -1;
    }

    return 0;
}
```

- lseek
- * Reads or write to a specific offset within a file
- * **Syntax:**

```
off_t lseek (int fd, off_t offset, int whence)
```

- fd - file descriptor
- offset - the offset of pointer within file (in bytes)
- whence - the method of offset

SEEK_SET - offset from the start of file (absolute)

SEEK_CUR - offset from current location + offset bytes (relative)

SEEK_END - offset from the end of file

- * Returns offset amount (in bytes) from the beginning of file
- * Returns -1 if error

Example

System Calls	Return Code	Current Offset
<code>fd = open("file", O_RDONLY);</code>	3	0
<code>lseek(fd, 200, SEEK_SET);</code>	200	200
<code>read(fd, buffer, 50);</code>	50	250
<code>close(fd);</code>	0	-

- rename (update file name)
 - * Is a system call
 - * Changes the name of file
 - * Is **atomic** (after crash, it will be either old or new, but not in-between)
 - * **Syntax:** `int rename(const char *old, const char *new)`
 - old - name of old file
 - new - name of new file
 - * Returns 0 if successful
 - * Returns -1 if error

Example

```

int fd = open("foo.txt.tmp", O_WRONLY|O_CREAT|O_TRUNC,
              S_IRUSR|S_IWUSR);
write(fd, buffer, size); // write out new version of file
fsync(fd);
close(fd);
rename("foo.txt.tmp", "foo.txt");

```

- stat (get file info)
 - * displays metadata of a certain file stored in **inode**
 - * **Syntax:** `int stat(const char *path, struct stat *buf)`
 - path - file descriptor of file that's being inquired
 - buf - A stat structure where data about the file will be stored (see below)

```

struct stat {
    dev_t      st_dev;        // ID of device containing file
    ino_t      st_ino;        // inode number
    mode_t     st_mode;       // protection
    nlink_t    st_nlink;      // number of hard links
    uid_t      st_uid;        // user ID of owner
    gid_t      st_gid;        // group ID of owner
    dev_t      st_rdev;       // device ID (if special file)
    off_t      st_size;       // total size, in bytes
    blksize_t  st_blksize;    // blocksize for filesystem I/O
    blkcnt_t   st_blocks;     // number of blocks allocated
    time_t     st_atime;      // time of last access
    time_t     st_mtime;      // time of last modification
    time_t     st_ctime;      // time of last status change
};

```

Figure 39.5: The stat structure.

Example

```

#include <unistd.h>
#include <stdio.h>
#include <sys/stat.h>
#include <sys/types.h>

int main(int argc, char **argv)
{
    if(argc != 2)
        return 1;

    struct stat fileStat;
    if(stat(argv[1], &fileStat) < 0)
        return 1;

    printf("Information for %s\n", argv[1]);
    printf("-----\n");
    printf("File Size: \t\t%d bytes\n", fileStat.st_size);
    printf("Number of Links: \t%d\n", fileStat.st_nlink);
    printf("File inode: \t\t%d\n", fileStat.st_ino);

    printf("File Permissions: \t");
    printf( (S_ISDIR(fileStat.st_mode)) ? "d" : "-");
    printf( (fileStat.st_mode & S_IRUSR) ? "r" : "-");
    printf( (fileStat.st_mode & S_IWUSR) ? "w" : "-");
    printf( (fileStat.st_mode & S_IXUSR) ? "x" : "-");
    printf( (fileStat.st_mode & S_IRGRP) ? "r" : "-");
    printf( (fileStat.st_mode & S_IWGRP) ? "w" : "-");
    printf( (fileStat.st_mode & S_IXGRP) ? "x" : "-");
    printf( (fileStat.st_mode & S_IROTH) ? "r" : "-");
    printf( (fileStat.st_mode & S_IWOTH) ? "w" : "-");
    printf( (fileStat.st_mode & S_IXOTH) ? "x" : "-");
    printf("\n\n");

    printf("The file %s a symbolic link\n", (S_ISLNK(fileStat.st_mode)) ? "is" : "is not");

    return 0;
}

```

The result of above is:

```
$ ./testProgram testfile.sh

Information for testfile.sh
-----
File Size:          36 bytes
Number of Links:    1
File inode:         180055
File Permissions:   -rwxr-xr-x

The file is not a symbolic link
```

- `unlink` (removing file)
 - Is a system call
 - Removes a file (including symbolic link) from the system
 - **Syntax:** `int unlink(const char *pathname)`
 - * `pathname` - path to file
 - Returns 0 if successful
 - Returns -1 if error

Example

```
#include <unistd.h>

char *path = "/modules/pass1";
int  status;
...
status = unlink(path);
```

- `mkdir` (creating directory)
 - Is a system call
 - **Syntax:** `int mkdir(const char *path, mode_t mode)`
 - * `path` - path of directory (including name)
 - * `mode` - permission group
 - Returns 0 if successful
 - Returns -1 if error
 - directories can never be written directly
 - * directory is in format called **File System Metadata**
 - * directory can only be updated directly
 - creates two directories on creation `.` (current) and `..` (parent)

Example


```
#include <sys/types.h>
#include <sys/stat.h>

int status;
...
status = mkdir("/home/cnd/mod1", S_IRWXU | S_IRWXG | S_IROTH | S_IXOTH);
```

- `opendir`, `readdir`, `closedir` (reading directory)
 - Are system calls
 - Are under `<dirent.h>` library
 - Requires struct `dirent` data structure

```
struct dirent {
    char          d_name[256]; // filename
    ino_t          d_ino;       // inode number
    off_t          d_off;       // offset to the next dirent
    unsigned short d_reclen;    // length of this record
    unsigned char  d_type;      // type of file
};
```

- **Syntax (`opendir`):** `DIR *opendir(const char *dirname)`
 - * `dirname` - directory path
 - * Returns a pointer to the directory stream
 - * The stream is positioned at the first entry in the directory.
- **Syntax (`readdir`):** `struct dirent *readdir(DIR *dirp);`
 - * `dirp` - directory stream
 - * Returns a pointer to a `dirent` structure representing the next directory entry in the directory stream
 - * Returns `NULL` on reaching the end of the directory stream
- **Syntax (`closedir`):** `int closedir(DIR *dirp);`
 - * `dirp` - directory stream
 - * Returns 0 if successful
 - * Returns -1 otherwise

Example

```
int main(int argc, char *argv[]) {
    DIR *dp = opendir(".");
    assert(dp != NULL);
    struct dirent *d;
    while ((d = readdir(dp)) != NULL) {
        printf("%lu %s\n", (unsigned long) d->d_ino,
              d->d_name);
    }
    closedir(dp);
    return 0;
}
```

- **rmdir** (Deleting Directories)
 - Removes a directory whose name is given by path
 - Is performed only when directory is empty
 - Is included in `<unistd.h>` library
 - Fails if is symbolic link
 - **Syntax:** `int rmdir(const char *path)`
 - * path - path of directory
 - Returns 0 if successful
 - Returns -1 if error

Example

```
#include <unistd.h>

int status;
...
status = rmdir("/home/cnd/mod1");
```

- **unlink** (Remove file)
 - Remove a link to a file
 - Is called **unlink** because it decrements **reference count** in inode
 - * Deletes file completely when reference count within the inode number is 0
 - **Syntax:**

```
#include <unistd.h>

int unlink(const char *pathname);
```

 - * pathname - pathname to file
 - Returns 0 if successful
 - Returns -1 if error
 - Is used by linux command `rm`

Example

```
#include <unistd.h>
```

```
char *path = "/modules/pass1";
int  status;
...
status = unlink(path);
```

```
prompt> echo hello > file
prompt> stat file
... Inode: 67158084    Links: 1 ...
prompt> ln file file2
prompt> stat file
... Inode: 67158084    Links: 2 ...
prompt> stat file2
... Inode: 67158084    Links: 2 ...
prompt> ln file2 file3
prompt> stat file
... Inode: 67158084    Links: 3 ...
prompt> rm file
prompt> stat file2
... Inode: 67158084    Links: 2 ...
prompt> rm file2
prompt> stat file3
... Inode: 67158084    Links: 1 ...
prompt> rm file3
```

- **Symbolic Link:**

- Is directory entry containing "true" path to the file
- Is a shortcut that reference to a file instead of inode value ^[2]



2



3



- **Hard Link:**

- Is a direct reference to a file via its inode [2]
- Is second directory entry identical to first

1



2



3



4

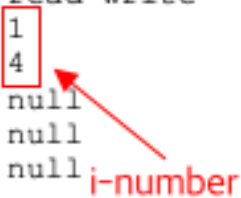


- **Crash Consistency**
 - Inode before update

```

owner      : remzi
permissions : read-write
size       : 1
pointer    : 4
pointer    : null
pointer    : null
pointer    : null

```

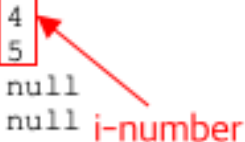


– Inode after update

```

owner      : remzi
permissions : read-write
size       : 2
pointer    : 4
pointer    : 5
pointer    : null
pointer    : null

```



References

- 1) codewiki, stat, link
- 2) The Open Group Base Specification, unlink, link


e) a) Notes

- What is a sector? What is a sector address?
- How can I get to specific inode from block (e.g inode # 32 on block 2)?
- **Index Based File System**
 - Disadvantages
 1. Uses a large amount of metadata per file (especially for large files)

Example

Linux's ext2, ext3

- **Extent Based File System**



- Requires only a disk pointer + length (in blocks)
- Is also called **contiguous allocation**
- Is simple
- Is less flexible but more compact
- Works well when there is enough free space on the disk and files can be laid out contiguously

Example

Linux's ext 4

- **inode**
 - Inode block computation

$$\text{block number} = (\text{inode \#} * \text{sizeof}(\text{inode})) / \text{block size} \quad (1)$$

Example

Target: inode #32

Inode Size: 256 bytes

Block Size: 4096 bytes

$$\text{block number} = (\text{inode \#} * \text{sizeof}(\text{inode})) / \text{block size} \quad (2)$$

$$= \frac{32 * 256}{4096} \quad (3)$$

$$= 2 \quad (4)$$

- **superblock**

- Contains information about the following
 - * The number of inodes and data blocks in a particular file system
 - * The magic number of some kind to identify the file system type
 - * Where the inode table begins
- Is read first on mount before attaching to file system
- **inode/data bitmap**
 - uses bit to indicate whether the corresponding object/block is free
 - * 0 means free
 - * 1 means in use