1. a) 1) 4 - inode blocks. 1 for the file `c`, and 3 for the directdories `/`, `a`, `b`
       2) 3 - directory blocks - one for root `/`, one for `a`, the other for `b`
       3) 1 - single indirect block as far as we know. The file definitely has more than 12 blocks (# of data blocks pointed by direct pounters), but less than 1036 (# of data blocks pointed by direct pointers and single indirect pointers). We are reading block 1034.
       4) 1 - data block for file `c`

2. (a) All of the above

   Notes

   - **Inode**


```
owner        : remzi
permissions  : read-write
size         : 1
pointer      : 4
pointer      : null
pointer      : null
pointer      : null
```

       – Is short form of **index node**
       – Describes a file system object such as file or data
       – Contains all information about a file/directory, including
         * File Type,
         * Size
         * Number of blocks allocated to it
         * Protection information
         * Time information (e.g time created, time modified)
         * Location of data blocks residing on disk

   References

       1) Wikipedia, Inode, link
       2) Machanick, Philip. (2016). Teaching Operating Systems: Just Enough Abstraction. 642. 10.1007/978-3-319-47680-3_10., link

   (b) Size, the location of data blocks that reside on disk

   Notes

- I wonder what information about blocks inode has. Is it total number of blocks both inode and data, or just data?
- I struggled a bit on this one. I should find an easier way to remember which information inode has

(c)   • **Inode Bitmap and Data Block Bitmap**

```
(b) - Data Leak
(c) - Inode Leak
```
- **New Directory Inode**

```
(a) - No inconsistency
```
- **Inode Bitmap**, **Data Block Bitmap**, **Existing Directory Data**, **New Directory Inode**, and **New Directory Data**

```
(e) - Inconsistent inode data
```
- **Inode Bitmap**, and **New Directory Inode**

```
(c) - Inode leak
(d) - Multiple file paths may point to same inode
```
- **New Directory Inode**, **Existing Directory Inode**, **Existing Directory Data**

```
(e) - Inconsistent inode data
(f) - Something points to garbage
```

Correct Solution

- **Inode Bitmap and Data Block Bitmap**

```
(b) - Data Leak
(c) - Inode Leak
```
- **New Directory Inode**

```
(a) - No inconsistency
```
- **Inode Bitmap**, **Data Block Bitmap**, **Existing Directory Data**, **New Directory Inode**, and **New Directory Data**

```
(e) - Inconsistent inode data
```
- **Inode Bitmap**, and **New Directory Inode**

```
(c) - Inode leak
```
- **New Directory Inode**, **Existing Directory Inode**, **Existing Directory Data**

```
      (d) – Multiple file paths may point to same inode
      (f) – Something points to garbage
```

Notes

- I wonder how system call for reading file/directory works in UNIX. Does it check for bitmap?
- I wonder how system call for deleting file/directory works in UNIX
- I wonder how system call for creatubg file/directory works in UNIX
- Learned that
  - Missing Inode Bitmap - multiple file paths may point to same inode
- **File API**
  - open (create/access file)
    * Is a system call
    * Reads target inode into memory (when loading)
    * Does three things on creation
      1) make structure (inode) that racks all relevant information about file
      2) link human readible name to the file, and put that link to a directory
      3) increment **reference count** in inode
    * **Syntax:**

      ```
      int fd = open("foo". O_CREAT|O_WRONLY|O_TRUNC, S_IRUSR|S_IWUSR)
      ```

      · O_CREAT - Creates file "foo" if does not exist
      · O_WRONLY - Open file for writing only (default)
      · O_TRUNC - Overwrites existing file Need example/Clarification
      · Can have multiple flags
    * Returns **file descriptor** or fd for short
      · Is an integer
      · Is used to access a file
      · Is <u>private</u> per process
      · Can be used to read() and write() files

    ### Example

```
#include <fcntl.h>
...
int fd;
mode_t mode = S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH;
char *filename = "/tmp/file";
...
fd = open(filename, O_WRONLY | O_CREAT | O_TRUNC, mode);
...
```

File can be read by owner
File can also be written by owner
File can also be read by group
File can also be read by others

Means
1. File is Writable AND
2. Create file if doesn't exist AND
3. Overwrite file if exists

* Amount of I/O generated by `open()` is proportional to length of pathname (wait. How is I/O involved in open()?)

– (read) (read file)
   * Is a system call
   * **Syntax:**

   ssize_t read (int fd, void *buf, size_t count)

   · `fd` - file descriptor (from `open()`)
   · `buf` - container for the read data
   · `count` - number of bytes to read
   * Returns number of bytes read, if successful
   * Returns 0 if is at, or past the end of file

   <u>**Example**</u>

```
char buf[4096];
int fd = open("/a/b/c", 0); // open in read-only mode
lseek(fd, 1034*4096, 0);    // seek to position (1034*4096) from start of file
read(fd, buf, 4096);        // read 4k of data from file
```

System Calls	Return Code	Current Offset	
fd = open("file", O_RDONLY);	3	0	
read(fd, buffer, 100);	100	100	read continues
read(fd, buffer, 100);	100	200	for each call
read(fd, buffer, 100);	100	300	
read(fd, buffer, 100);	0	300	returns 0
close(fd);	0	–	if at end

4

– `write` (write file)
  * Is a system call
  * Writes data out of a buffer
  * **Syntax:**

    `ssize_t write (int fd, const void * buf, size_t nbytes)`

    · `fd` - file descriptor
    · `buf` - A pointer to a buffer to write to file
    · `nbytes` - number of bytes to write. If smaller than buffer, the output is truncated

Example

```c
#include <unistd.h>
#include <fcntl.h>

int main(void)
{
    int filedesc = open("testfile.txt", O_WRONLY | O_APPEND);

    if (filedesc < 0) {
        return -1;
    }

    if (write(filedesc, "This will be output to testfile.txt\n", 36) != 36) {
        write(2, "There was an error writing to testfile.txt\n", 43);
        return -1;
    }

    return 0;
}
```

– `lseek`
  * Reads or write to a specific offset within a file
  * **Syntax:**

    `off_t lseek (int fd, off_t offset, int whence)`

    · `fd` - file descriptor
    · `offset` - the offset of pointer within file (in bytes)
    · `whence` - the method of offset

        `SEEK_SET` - offset from the start of file (absolute)
        `SEEK_CUR` - offset from current location + offset bytes (relative)
        `SEEK_END` - offset from the end of file
  * Returns offset amount (in bytes) from the <u>beginning</u> of file
  * Returns -1 if error

Example

System Calls	Return Code	Current Offset
fd = open("file", O_RDONLY);	3	0
lseek(fd, 200, SEEK_SET);	200	200
read(fd, buffer, 50);	50	250
close(fd);	0	–

move 200 bytes from the start of file

read 50 bytes

- – rename (update file name)
  - ∗ Is a system call
  - ∗ Changes the name of file
  - ∗ Is **atomic** (after crash, it will be either old or new, but not in-between)
  - ∗ **Syntax:** `int rename(const char *old, const char *new)`
    - · `old` - name of old file
    - · `new` - name of new file
  - ∗ Returns 0 if successful
  - ∗ Returns -1 if error

Example

```
int fd = open("foo.txt.tmp", O_WRONLY|O_CREAT|O_TRUNC,
              S_IRUSR|S_IWUSR);
write(fd, buffer, size); // write out new version of file
fsync(fd);
close(fd);
rename("foo.txt.tmp", "foo.txt");
```

e.g. "hello\n"

- – stat (get file info)
  - ∗ displays metadata of a certain file stored in **inode**
  - ∗ **Syntax:** `int stat(const char *path, struct stat *buf)`
    - · `path` - file descriptor of file that's being inquired
    - · `buf` - A `stat` structure where data about the file will be stored (see below)

```
struct stat {
    dev_t     st_dev;      // ID of device containing file
    ino_t     st_ino;      // inode number
    mode_t    st_mode;     // protection
    nlink_t   st_nlink;    // number of hard links
    uid_t     st_uid;      // user ID of owner
    gid_t     st_gid;      // group ID of owner
    dev_t     st_rdev;     // device ID (if special file)
    off_t     st_size;     // total size, in bytes
    blksize_t st_blksize;  // blocksize for filesystem I/O
    blkcnt_t  st_blocks;   // number of blocks allocated
    time_t    st_atime;    // time of last access
    time_t    st_mtime;    // time of last modification
    time_t    st_ctime;    // time of last status change
};
```

Figure 39.5: **The stat structure.**

Example

```c
#include <unistd.h>
#include <stdio.h>
#include <sys/stat.h>
#include <sys/types.h>

int main(int argc, char **argv)
{
    if(argc != 2)
        return 1;

    struct stat fileStat;
    if(stat(argv[1],&fileStat) < 0)
        return 1;

    printf("Information for %s\n",argv[1]);
    printf("---------------------------\n");
    printf("File Size: \t\t%d bytes\n",fileStat.st_size);
    printf("Number of Links: \t%d\n",fileStat.st_nlink);
    printf("File inode: \t\t%d\n",fileStat.st_ino);

    printf("File Permissions: \t");
    printf( (S_ISDIR(fileStat.st_mode)) ? "d" : "-");
    printf( (fileStat.st_mode & S_IRUSR) ? "r" : "-");
    printf( (fileStat.st_mode & S_IWUSR) ? "w" : "-");
    printf( (fileStat.st_mode & S_IXUSR) ? "x" : "-");
    printf( (fileStat.st_mode & S_IRGRP) ? "r" : "-");
    printf( (fileStat.st_mode & S_IWGRP) ? "w" : "-");
    printf( (fileStat.st_mode & S_IXGRP) ? "x" : "-");
    printf( (fileStat.st_mode & S_IROTH) ? "r" : "-");
    printf( (fileStat.st_mode & S_IWOTH) ? "w" : "-");
    printf( (fileStat.st_mode & S_IXOTH) ? "x" : "-");
    printf("\n\n");

    printf("The file %s a symbolic link\n", (S_ISLNK(fileStat.st_mode)) ? "is" : "is not");

    return 0;
}
```

The result of above is:

```
$ ./testProgram testfile.sh

Information for testfile.sh
---------------------------
File Size:              36 bytes
Number of Links:        1
File inode:             180055
File Permissions:       -rwxr-xr-x

The file is not a symbolic link
```

- unlink (removing file)
  - Is a system call
  - Removes a <u>file</u> (including symbolic link) from the system
  - **Syntax:** int unlink(const char *pathname)
    * pathname - path to file
  - Returns 0 if successful
  - Returns -1 if error

Example

```
#include <unistd.h>


char *path = "/modules/pass1";
int    status;
...
status = unlink(path);
```

- `mkdir` (creating directory)
    - Is a system call
    - **Syntax:** `int mkdir(const char *path, mode_t mode)`
        * `path` - path of directory (including name)
        * `mode` - permission group
    - Returns 0 if successful
    - Returns -1 if error
    - directories can never be written directly
        * directory is in format called **File System Metadata**
        * directory can <u>only</u> be updated directly
    - creates two directories on creation `.` (current) and `..` (parent)

<u>Example</u>

```
#include <sys/types.h>
#include <sys/stat.h>


int status;
...
status = mkdir("/home/cnd/mod1", S_IRWXU | S_IRWXG | S_IROTH | S_IXOTH);
```

- `opendir, readdir, closedir` (reading directory)
    - Are system calls
    - Are under `<dirent.h>` library
    - Requires `struct dirent` data structure

```
struct dirent {
  char           d_name[256]; // filename
  ino_t          d_ino;       // inode number
  off_t          d_off;       // offset to the next dirent
  unsigned short d_reclen;    // length of this record
  unsigned char  d_type;      // type of file
};
```

    - **Syntax (opendir):** `DIR *opendir(const char *dirname)`

* dirname - directory path
* Returns a pointer to the directory stream
* The stream is positioned at the first entry in the directory.
- **Syntax (readdir):** `struct dirent *readdir(DIR *dirp);`
  * `dirp` - directory stream
  * Returns a pointer to a dirent structure representing the next directory entry in the directory stream
  * Returns NULL on reaching the end of the directory stream
- **Syntax (closedir):** `int closedir(DIR *dirp));`
  * `dirp` - directory stream
  * Returns 0 if successful
  * Returns -1 otherwise

Example

```
int main(int argc, char *argv[]) {
    DIR *dp = opendir(".");
    assert(dp != NULL);
    struct dirent *d;
    while ((d = readdir(dp)) != NULL) {
        printf("%lu %s\n", (unsigned long) d->d_ino,
                d->d_name);
    }
    closedir(dp);
    return 0;
}
```

* rmdir (Deleting Directories)
  - Removes a directory whose name is given by path
  - Is performed only when directory is empty
  - Is included in <unistd.h> library
  - Fails if is symbolic link
  - **Syntax:** `int rmdir(const char *path)`
    * `path` - path of directory
  - Returns 0 if successful
  - Returns -1 if error

Example

```
#include <unistd.h>


int status;
...
status = rmdir("/home/cnd/mod1");
```

- unlink (Remove file)
  - Remove a link to a file
  - Is called **unlink** because it decrements **reference count** in inode
    * Deletes file completely when reference count within the inode number is 0
  - **Syntax:**

    #include <unistd.h>

    int unlink(const char *pathname);

    * pathname - pathname to file
  - Returns 0 if successful
  - Returns -1 if error
  - Is used by linux command rm

Example

```
#include <unistd.h>


char *path = "/modules/pass1";
int   status;
...
status = unlink(path);
```

```
prompt> echo hello > file
prompt> stat file
... Inode: 67158084      Links: 1 ...
prompt> ln file file2
prompt> stat file
... Inode: 67158084      Links: 2 ...
prompt> stat file2
... Inode: 67158084      Links: 2 ...
prompt> ln file2 file3
prompt> stat file
... Inode: 67158084      Links: 3 ...
prompt> rm file
prompt> stat file2
... Inode: 67158084      Links: 2 ...
prompt> rm file2
prompt> stat file3
... Inode: 67158084      Links: 1 ...
prompt> rm file3
```

- **Symbolic Link:**
  - Is directory entry containing "true" path to the file
  - Is a shortcut that reference to a file instead of inode value [2]

2

3

- **Hard Link:**
  - Is a direct reference to a file via its inode [2]
  - Is second directory entry identical to first

1

- **Crash Consistency**
  - Inode before update

```
owner        : remzi
permissions  : read-write
size         : 1
pointer      : 4
pointer      : null
pointer      : null
pointer      : null i-number
```

– Inode after update

```
owner        : remzi
permissions  : read-write
size         : 2
pointer      : 4
pointer      : 5
pointer      : null
pointer      : null i-number
```

References

1) codewiki, stat, link
2) The Open Group Base Specification, unlink, link

3. a) Indexed-based file system uses inode number and pointers to find data blocks, and data blocks can be set and anywhere, so external fragmentation isn't a problem

   b) Extent based file system only requires a pointer to first data block of file, and the rest is read by traveling contiguously, and this requires less disk block access than index-based file system of which has to go to inode block, indirect pointers and data blocks to go to a particular byte in file.

Notes

- What is a sector? What is a sector address?
- How can I get to specific inode from block (e.g inode # 32 on block 2)?
- I should record differences between linked-list-based FS, Extent-based FS, and indexed FS
- **Index Based File System**

- Has 15 blocks of pointers that points to either inode, indirect pointers, or data block
- No external fragmentation
- Files can be easily grown

Example

Linux's ext2, ext3

- **Extent Based File System**

- Requires only a disk pointer + length (in blocks)
- Is also called **contiguous allocation**
- Is simple

- Is less flexible but more compact
- Works well when there is enough free space on the disk and files can be laid out contiguously

  Example

  Linux's ext 4

- **inode**
  - Inode block computation

  $$\texttt{block number} = (\texttt{inode \# } * \texttt{ sizeof(inode)})/\texttt{block size} \quad (1)$$

  Example

  Target: inode #32
  Inode Size: 256 bytes
  Block Size: 4096 bytes

  $$\texttt{block number} = (\texttt{inode \# } * \texttt{ sizeof(inode)})/\texttt{block size} \quad (2)$$
  $$= \frac{32 * 256}{4096} \quad (3)$$
  $$= 2 \quad (4)$$

- **superblock**
  - Contains information about the following
    * The number of inodes and data blocks in a particular file system
    * The magic number of some knd to identify the file system type
    * Where the inode table begins
  - Is read first on mount before attaching to file system

- **inode/data bitmap**
  - Accessed only when allocation/deallocation is needed
    * `Read()` $\rightarrow$ no bitmap required
  - Uses bit to indicate whether the corres object/block is free
    * 0 means free
    * 1 means in use

- **Reading a File from Disk**
  Example

  When

```
open("/foo/bar", O_READONLY)
```

is called

- the goal is to find the inode of the file `bar` to read its basic information (i.e. includes permission, information, file size etc)
- done by traversing the pathname and locate the desired inode
- Steps
    1. Find **inode** of the root directory by looking for **i-number** (or **inode number**)
        * Root directory has no parent directory
        * Root directory's **inode number** is 2 (for UNIX file systems)
    2. Read the **inode** of root directory
    3. Once its **inode** is read, read through its directory data (pointers to **data blocks**) until the inode number of `foo` is found (e.g 42)
    4. Recursively traverse the pathname until the desired inode is found (more specifically, the **inode number** of bar)
    5. Issue a `open()` to read `bar`'s inode to memory
    6. Issue a `read()` system call to read from file `bar`
        * without `lseek()`, reads file from the first file data block (e.g. `bar data[0]`)
        * `lseek(..., offset_amt * size_of_file_block)` is used to offset/move to desired block in `bar`
    7. Trasnfer data to `buf` data block
    8. Read until `read()` returns 0, or desired data block has been read
    9. Close `fd`. No I/O is read.

- **Writing to Disk**

Given a call

`create(...)` (Note: `open` to be exact)

- 5 I/Os are generated per write
  * Read inode (to traverse to the location of new data block)
  * Reading data bitmap
  * Writing data bitmap
  * Write data block
  * Write inode (to update data block's location in inode)
- 10 I/Os are generated per file creation:
  * Read inode bitmap (to find free inode)
  * Write inode bitmap (to mark it allocated)
  * Create one new inode (to initialize it)
  * Write the location of new inode block in `foo` (by linking high-level name of file `bar` to its inode number and storing in data block)
  * Perform one read and write to the directory inode and update it

- **Static Partitioning**
  - Divides resources into fixed proportion <u>once</u>
    * e.g. two possible users of memory → give fraction of memory to one user and rest to the other

- Advantages
  * Ensures each user receives some share of the resource
  * Delivers more predictable performance (usually)
  * Easier to implement
- Disadvantages
  * Is wasteful
  *

- **Dynamic Partitioning**
  - Gives out different amounts of resources over time
  - Lets resource-hungry users consume idle resources
  - Advantages
    * Flexible
    * Can achieve better utilization than **static partitioning**
  - Disadvantages
    * More complex to implement
    * Could lead to worse performance
      · e.g idle resource got consumed by others and take long time to reclaim it when needed (the perodic frozen feeling when loading screen)

References

1) Columbia University, Operating Systems, link

4. 1) i) all of the above
    ii) To minimize damage, it should be updated in the order of

    1. Data Region
    2. Inode Table
    3. Data Bitmap
    4. Inode Bitmap

    The reason is that when data block and/or inode block are set and crash happens, the file system treats as if nothing had happened.

    And the reason is that before data/inode block is allocated, it first checks the inode and data bitmap, which contains information about whether the block is occupied or not.

    Once they are allocated, we risk data/inode leak, and if done improperly, the data/inode block would not be available until actions are taken.

2)   i) inode bitmap, data bitmap

    ii) To minimize damage, it should be updated in the order of

      1. Data Bitmap

      2. Inode Bitmap

    The reason is that when data bitmap is removed and crash happens, The data are still in place, and those can be used to continue the operation (assuming that nothing is disturbed).

    More specifically, deletion of file is complete when reference count in the inode of file B hits 0, and as long as data is in place we can work so that the inode's reference count hits 0, and the file is removed.

    However, if done the opposite and crash occurs, we can't remove its data bitmaps, and data leak would result.

    ### Notes
      • I really need to know what happens

3) The primary motiviations behind log-structured file system are as follows

    1. Decrease a large number of intermediary processes during a write in FFS (e.g creation of data block, inode)

    2. Utilize the ever increasing capacity of memory, and write a large block of information at once (to boost performance)

    3. To compensate for the slow development of cheap and fast motor for ever increasing capacity and the need for higher data transfer rate

The update process for log-structured file system works by storing write information to in-memory space called **segment**, and write all once once its at capacity.

On the other hand, the writing process of FFS works by traveling to target directory, create a inode bitmap, create file inode, update the currect directory inode, create a data block in directory inode that points to the file inode, allocate data inode, and write data block for file, and update the current file meta data.

And this data block writing process continues until all are filled.

Notes

• **Log Structured File System**
    – Wait. This sounds very similar to **extent-based file system**
    – Buffers all updates (including metadata) in an in-memory **segment**, and when segment is full, it is written to disk in one long, sequential transfer to unused part of the disk

- Instead of overwriting files, always writes unused portion of the disk, and reclaim the old space through cleaning
- Motivations
  1. **System memories are growing**
     * Data is cached in memory
     * Reads are serviced by cache
     * Disk traffic is increasingly consists of writes
     * File performance $\approx$ write performance
  2. **There is a large gap between random I/O performance and sequential I/O performance**
     * More bits stored on hard drive $\Rightarrow$ bandwith of acessing bits $\uparrow$
     * Harder to create cheap, small motors that spin platters faster, and move arm more quickly
  3. **Existing file systems perform poorly on many common workloads**
     * Many intermediary writes performed per data block (e.g. Bitmap, inode, data block)
     * Many short seeks + rotation delays = performance less than the peak
  4. **File systems are not raid aware**
- How it works (Writing to Disk)

  Basic idea: Write all updates (e.g. data blocks inodes) to the disk sequentially (**write buffering**)

  1. Buffer updates in an in-memory **segment**
  2. Write the **segment** all at once sequentially when received sufficient number of updates
- Advatages
  1. Has very high performance
- Disadvantages
  1. Is complex
  2. Generates lots of garbages
  3. Scattered old data. Needs to run **compaction** periodically [2]
- **Fast File System**
  - Divides inode tables into chunks and stores in different cylinder groups
  - Advantages
    * No external fragmentation
  - Disadvantages
    * Extra overhead: creates and updates many intermediary files (inode, data block) during a write

Refernces:

1) Ousterhout J. (1991). *The Design and Implementation of a Log-Structured File System*. link

2) Cornell University, Log-Structured File Systems, link

4) The challenge of locating data and metadata on LFS is that both data block and inodes are not at the fixed position.

To compensate for this, LFS has imap that contains the pointers to (i.e. address of) the inode with the latest information, but imap is also placed along side inode and data block.

To compensate for this, LFS has checkpoint region at fixed location that contains the pointers to (i.e. address of) the imap with the latest information.

So when LFS loads, it first reads the checkpoint regions, upload all imap to in-memory, and when file system gives target inode number, it finds inode and its data blocks from there much like UNIX file system.

Notes

- **Log Structured File System (cont')**
  - Data Block
    * Is found by looking for **inode** that point to it

  - Inode
    * Stores metadata of file
    * Is the same as UNIX file system inode
      · Has indirect pointers
      · Has direct pointers
      · Has double indirect pointers
    * Is found through a data structure called **inode map**
    * Are scattered throughout the disk & keeps moving after update

Placed all togheter onto the disk

- Inode Map
  * Is a data structure
  * Is also moving
    · new data block, inode and imap are placed contiguously
  * Locates inode
  * Takes inode number as input
  * Produces the disk address of the most recent version of the inode
- Checkpoint Region
  * Is fixed
  * Locates imap
  * Contains pointers to (i.e. address of) the latest pieces of inode map
  * Is updated periodically (e.g every 30 seconds)
    · Done to ill-affect performance
- How it works (Reading the Disk)
  1. Read the checkpoint region
  2. Read the entire inode map and cache it in memory
  3. Locate inode given inode number of file
  4. Proceed the same as typical UNIX file system

5. a) **Notes**

- **Crash Consistency Problem**
  - Desired: **atomic** updates. That is, on crash, the file on write is either in (state 1 - before the file got updated) or (state 2 - after the file got updated)
  - Reality: This is not possible
  - Is the reason why computers have 'Don't turn off computer' message
- **File System Checker**
  - Is implemented in early file system
  - Is used by UNIX tool **fsck** ('file system checker')
  - Basic Idea: Let inconsistencies happen and fix them later (when rebooting)

- Summary of how it works
  * **Inode State**
    · Corruption in file is checked (e.g. does it have valid file type such as directory file, or links)
    · If Inode cannot be fixed easily, it is removed, and its bitmap is updated
  * **Inode links**
  * **Duplicates**
  * **Bad Blocks**
  * **Directory Checks**
- Disadvantage
  * Way too slow. May take Hours.
  * Doesn't solve all problems (e.g. inode with incorrect data blocks)