

# 1 Exam Related Questions and Tips

- I wonder how system call for reading file/directory works in UNIX. Does it check for bitmap?
- I wonder how system call for deleting file/directory works in UNIX
- I wonder how system call for creatubg file/directory works in UNIX
- Learned that
  - Missing Inode Bitmap - multiple file paths may point to same inode

# 2 File API

- open (create/access file)
  - Is a system call
  - Reads target inode into memory (when loading)
  - Does three things on creation
    - 1) make structure (inode) that racks all relevant information about file
    - 2) link human readable name to the file, and put that link to a directory
    - 3) increment **reference count** in inode

– **Syntax:**

```
int fd = open("foo". O_CREAT|O_WRONLY|O_TRUNC, S_IRUSR|S_IWUSR)
```

- \* O\_CREAT - Creates file "foo" if does not exist
- \* O\_WRONLY - Open file for writing only (default)
- \* O\_TRUNC - Overwrites existing file **Need example/Clarification**
- \* Can have multiple flags

– Returns **file descriptor** or fd for short

- \* Is an integer
- \* Is used to access a file
- \* Is private per process
- \* Can be used to read() and write() files

## Example

```

#include <fcntl.h>
...
int fd;
mode_t mode = S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH;
char *filename = "/tmp/file";
...
fd = open(filename, O_WRONLY | O_CREAT | O_TRUNC, mode);
...

```

File can be read by owner  
 File can also be written by owner  
 File can also be read by group  
 File can also be read by others

Means

1. File is Writable AND
2. Create file if doesn't exist AND
3. Overwrite file if exists

- Amount of I/O generated by `open()` is proportional to length of pathname (wait. How is I/O involved in `open()`?)

- (read) (read file)

- Is a system call
- **Syntax:**

```
ssize_t read (int fd, void *buf, size_t count)
```

- \* `fd` - file descriptor (from `open()`)
- \* `buf` - container for the read data
- \* `count` - number of bytes to read

- Returns number of bytes read, if successful
- Returns 0 if is at, or past the end of file

### Example

```

char buf[4096];
int fd = open("/a/b/c", 0); // open in read-only mode
lseek(fd, 1034*4096, 0);    // seek to position (1034*4096) from start of file
read(fd, buf, 4096);        // read 4k of data from file

```

System Calls	Return Code	Current Offset	
fd = open("file", O_RDONLY);	3	0	
read(fd, buffer, 100);	100	100	← read continues for each call
read(fd, buffer, 100);	100	200	
read(fd, buffer, 100);	100	300	
read(fd, buffer, 100);	0	300	← returns 0 if at end
close(fd);	0	-	

- write (write file)

- Is a system call
- Writes data out of a buffer
- **Syntax:**

```
ssize_t write (int fd, const void * buf, size_t nbytes)
```

- \* fd - file descriptor
- \* buf - A pointer to a buffer to write to file
- \* nbytes - number of bytes to write. If smaller than buffer, the output is truncated

### Example

```
#include <unistd.h>
#include <fcntl.h>

int main(void)
{
    int filedesc = open("testfile.txt", O_WRONLY | O_APPEND);

    if (filedesc < 0) {
        return -1;
    }

    if (write(filedesc, "This will be output to testfile.txt\n", 36) != 36) {
        write(2, "There was an error writing to testfile.txt\n", 43);
        return -1;
    }

    return 0;
}
```

- lseek

- Reads or write to a specific offset within a file

– **Syntax:**

```
off_t lseek (int fd, off_t offset, int whence)
```

- \* fd - file descriptor
- \* offset - the offset of pointer within file (in bytes)
- \* whence - the method of offset

SEEK\_SET - offset from the start of file (absolute)

SEEK\_CUR - offset from current location + offset bytes (relative)

SEEK\_END - offset from the end of file

- Returns offset amount (in bytes) from the beginning of file
- Returns -1 if error

Example

System Calls	Return Code	Current Offset
fd = open("file", O_RDONLY);	3	0
lseek(fd, 200, SEEK_SET);	200	200
read(fd, buffer, 50);	50	250
close(fd);	0	-

move 200 bytes from the start of file

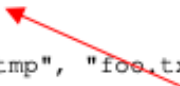
read 50 bytes

- rename (update file name)
  - Is a system call
  - Changes the name of file
  - Is **atomic** (after crash, it will be either old or new, but not in-between)
  - **Syntax:** `int rename(const char *old, const char *new)`
    - \* old - name of old file
    - \* new - name of new file
  - Returns 0 if successful
  - Returns -1 if error

Example

```
int fd = open("foo.txt.tmp", O_WRONLY|O_CREAT|O_TRUNC,
              S_IRUSR|S_IWUSR);
write(fd, buffer, size); // write out new version of file
fsync(fd);
close(fd);
rename("foo.txt.tmp", "foo.txt");
```

e.g. "hello\n"



- stat (get file info)
  - displays metadata of a certain file stored in **inode**
  - **Syntax:** `int stat(const char *path, struct stat *buf)`
    - \* path - file descriptor of file that's being inquired
    - \* buf - A stat structure where data about the file will be stored (see below)

```
struct stat {
    dev_t      st_dev;        // ID of device containing file
    ino_t      st_ino;        // inode number
    mode_t     st_mode;       // protection
    nlink_t    st_nlink;      // number of hard links
    uid_t      st_uid;        // user ID of owner
    gid_t      st_gid;        // group ID of owner
    dev_t      st_rdev;       // device ID (if special file)
    off_t      st_size;       // total size, in bytes
    blksize_t  st_blksize;    // blocksize for filesystem I/O
    blkcnt_t   st_blocks;     // number of blocks allocated
    time_t     st_atime;      // time of last access
    time_t     st_mtime;      // time of last modification
    time_t     st_ctime;      // time of last status change
};
```

Figure 39.5: The **stat** structure.

## Example

```

#include <unistd.h>
#include <stdio.h>
#include <sys/stat.h>
#include <sys/types.h>

int main(int argc, char **argv)
{
    if(argc != 2)
        return 1;

    struct stat fileStat;
    if(stat(argv[1], &fileStat) < 0)
        return 1;

    printf("Information for %s\n", argv[1]);
    printf("-----\n");
    printf("File Size: \t\t%d bytes\n", fileStat.st_size);
    printf("Number of Links: \t%d\n", fileStat.st_nlink);
    printf("File inode: \t\t%d\n", fileStat.st_ino);

    printf("File Permissions: \t");
    printf( (S_ISDIR(fileStat.st_mode)) ? "d" : "-");
    printf( (fileStat.st_mode & S_IRUSR) ? "r" : "-");
    printf( (fileStat.st_mode & S_IWUSR) ? "w" : "-");
    printf( (fileStat.st_mode & S_IXUSR) ? "x" : "-");
    printf( (fileStat.st_mode & S_IRGRP) ? "r" : "-");
    printf( (fileStat.st_mode & S_IWGRP) ? "w" : "-");
    printf( (fileStat.st_mode & S_IXGRP) ? "x" : "-");
    printf( (fileStat.st_mode & S_IROTH) ? "r" : "-");
    printf( (fileStat.st_mode & S_IWOTH) ? "w" : "-");
    printf( (fileStat.st_mode & S_IXOTH) ? "x" : "-");
    printf("\n\n");

    printf("The file %s a symbolic link\n", (S_ISLNK(fileStat.st_mode)) ? "is" : "is not");

    return 0;
}

```

The result of above is:

```

$ ./testProgram testfile.sh

Information for testfile.sh
-----
File Size:                36 bytes
Number of Links:          1
File inode:                180055
File Permissions:         -rwxr-xr-x

The file is not a symbolic link

```

- unlink (removing file)
  - Is a system call
  - Removes a file (including symbolic link) from the system
  - **Syntax:** int unlink(const char \*pathname)

- \* pathname - path to file
- Returns 0 if successful
- Returns -1 if error

### Example

```
#include <unistd.h>

char *path = "/modules/pass1";
int  status;
...
status = unlink(path);
```

- `mkdir` (creating directory)
  - Is a system call
  - **Syntax:** `int mkdir(const char *path, mode_t mode)`
    - \* path - path of directory (including name)
    - \* mode - permission group
  - Returns 0 if successful
  - Returns -1 if error
  - directories can never be written directly
    - \* directory is in format called **File System Metadata**
    - \* directory can only be updated directly
  - creates two directories on creation `.` (current) and `..` (parent)

### Example

```
#include <sys/types.h>
#include <sys/stat.h>

int status;
...
status = mkdir("/home/cnd/mod1", S_IRWXU | S_IRWXG | S_IROTH | S_IXOTH);
```

- `opendir`, `readdir`, `closedir` (reading directory)

- Are system calls
- Are under `<dirent.h>` library
- Requires `struct dirent` data structure

```
struct dirent {
    char      d_name[256]; // filename
    ino_t      d_ino;      // inode number
    off_t      d_off;      // offset to the next dirent
    unsigned short d_reclen; // length of this record
    unsigned char d_type;   // type of file
};
```

- **Syntax (`opendir`):** `DIR *opendir(const char *dirname)`
  - \* `dirname` - directory path
  - \* Returns a pointer to the directory stream
  - \* The stream is positioned at the first entry in the directory.
- **Syntax (`readdir`):** `struct dirent *readdir(DIR *dirp);`
  - \* `dirp` - directory stream
  - \* Returns a pointer to a `dirent` structure representing the next directory entry in the directory stream
  - \* Returns `NULL` on reaching the end of the directory stream
- **Syntax (`closedir`):** `int closedir(DIR *dirp);`
  - \* `dirp` - directory stream
  - \* Returns 0 if successful
  - \* Returns -1 otherwise

### Example

```
int main(int argc, char *argv[]) {
    DIR *dp = opendir(".");
    assert(dp != NULL);
    struct dirent *d;
    while ((d = readdir(dp)) != NULL) {
        printf("%lu %s\n", (unsigned long) d->d_ino,
              d->d_name);
    }
    closedir(dp);
    return 0;
}
```



- rmdir (Deleting Directories)
  - \* Removes a directory whose name is given by path
  - \* Is performed only when directory is empty
  - \* Is included in <unistd.h> library
  - \* Fails if is symbolic link
  - \* **Syntax:** `int rmdir(const char *path)`
    - path - path of directory
  - \* Returns 0 if successful
  - \* Returns -1 if error

### Example

```
#include <unistd.h>

int status;
...
status = rmdir("/home/cnd/mod1");
```

- unlink (Remove file)
  - \* Remove a link to a file
  - \* Is called **unlink** because it decrements **reference count** in inode
    - Deletes file completely when reference count within the inode number is 0
  - \* **Syntax:**  
  
`#include <unistd.h>`  
  
`int unlink(const char *pathname);`
    - pathname - pathname to file
  - \* Returns 0 if successful
  - \* Returns -1 if error
  - \* Is used by linux command rm

### Example

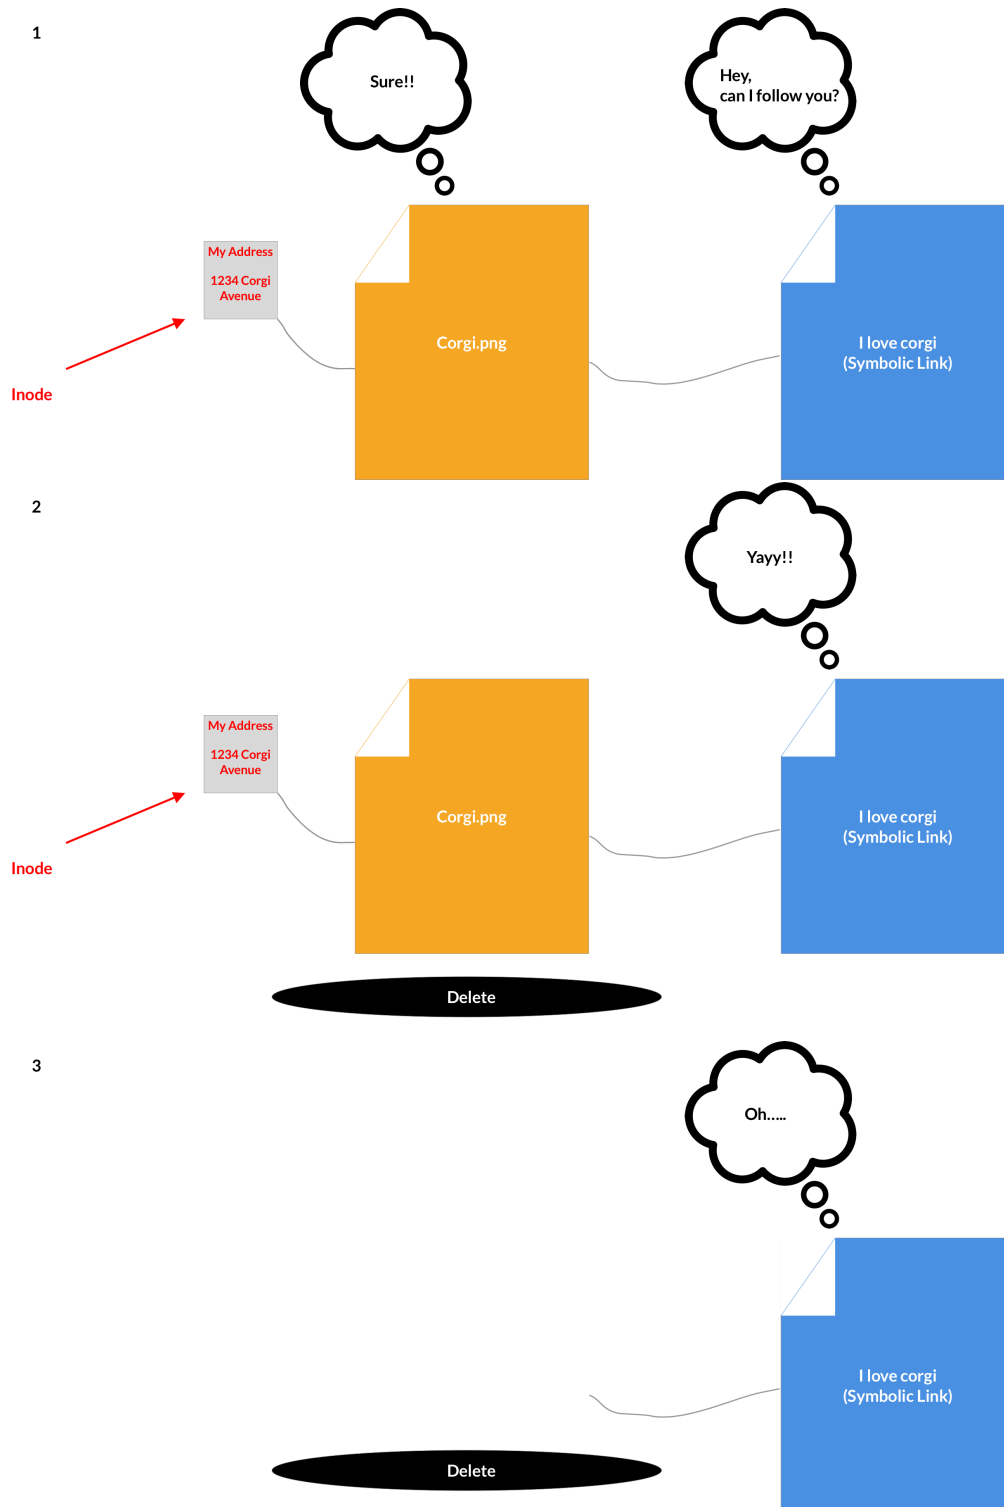
```
#include <unistd.h>
```

```
char *path = "/modules/pass1";  
int  status;  
...  
status = unlink(path);
```

```
prompt> echo hello > file  
prompt> stat file  
... Inode: 67158084    Links: 1 ...  
prompt> ln file file2  
prompt> stat file  
... Inode: 67158084    Links: 2 ...  
prompt> stat file2  
... Inode: 67158084    Links: 2 ...  
prompt> ln file2 file3  
prompt> stat file  
... Inode: 67158084    Links: 3 ...  
prompt> rm file  
prompt> stat file2  
... Inode: 67158084    Links: 2 ...  
prompt> rm file2  
prompt> stat file3  
... Inode: 67158084    Links: 1 ...  
prompt> rm file3
```

### 3 Symbolic Link:

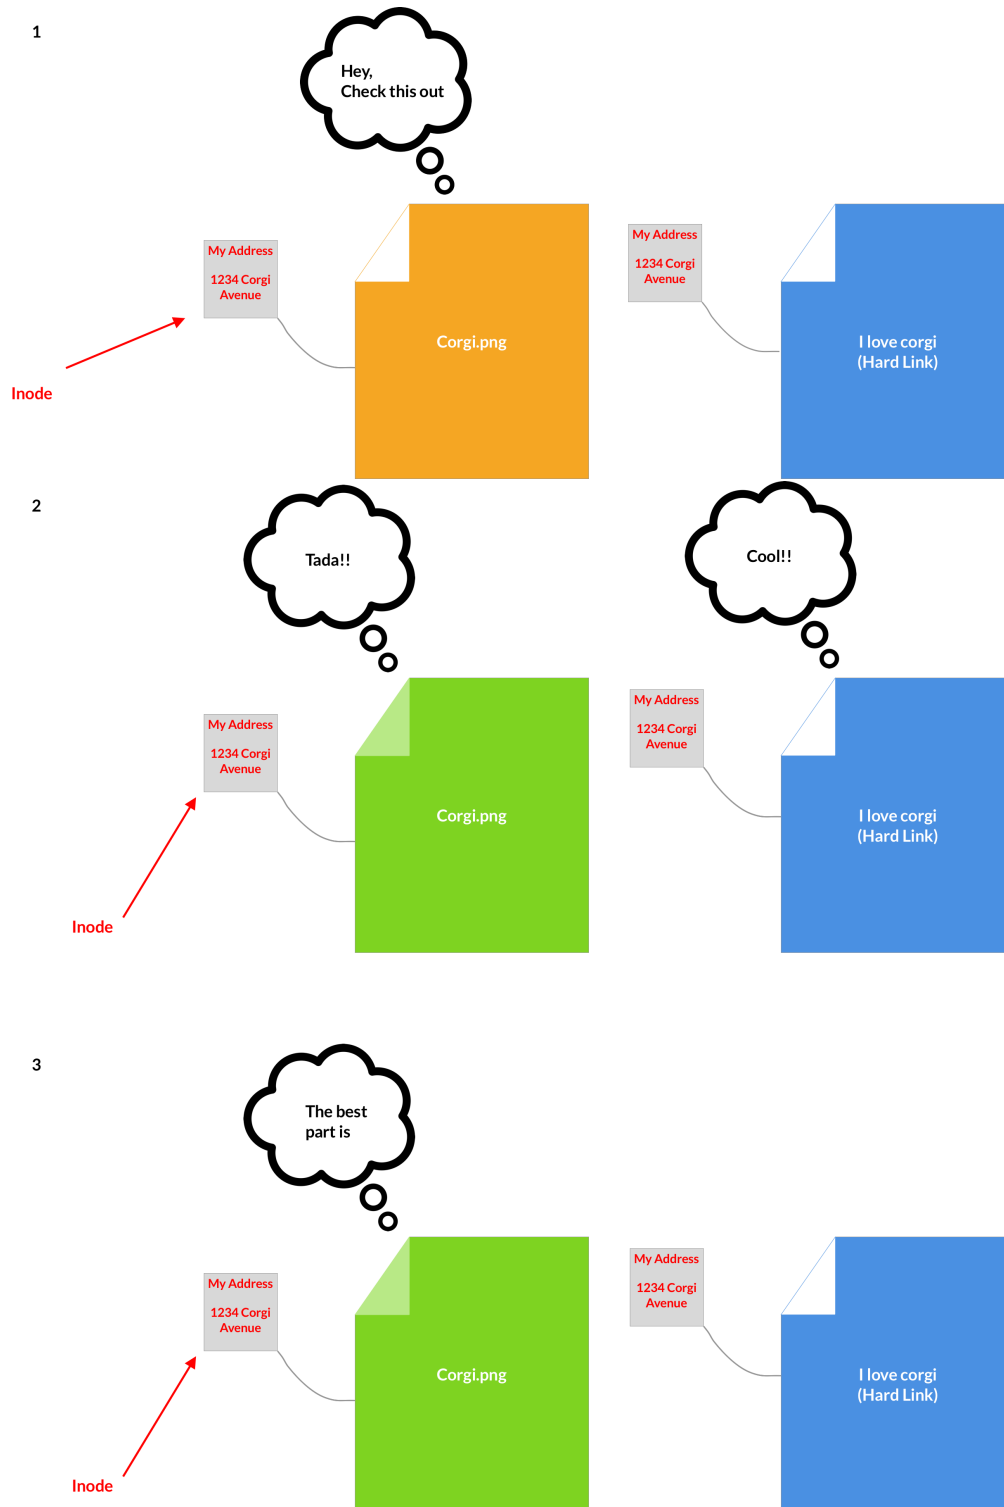
- Is directory entry containing "true" path to the file
- Is a shortcut that reference to a file instead of inode value <sup>[2]</sup>



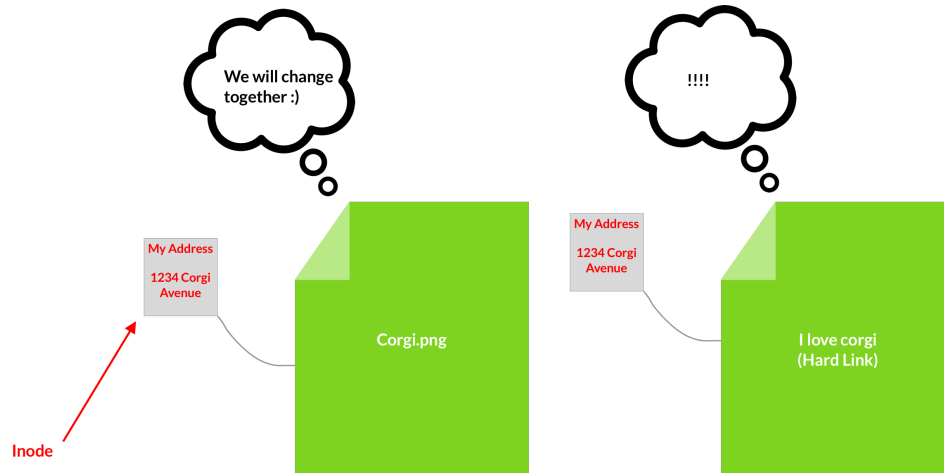
## 4 Hard Link:

- Is a direct reference to a file via its inode <sup>[2]</sup>

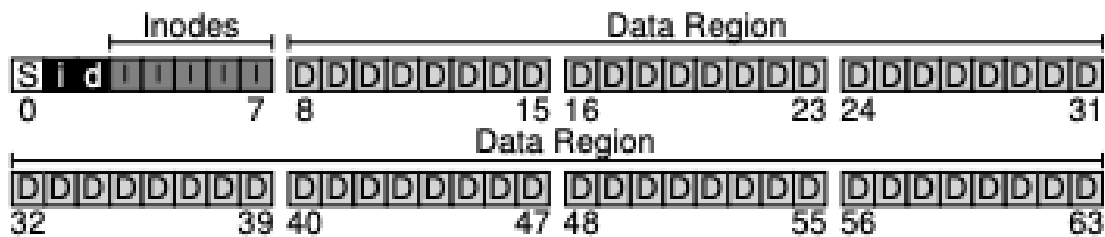
- Is second directory entry identical to first



4

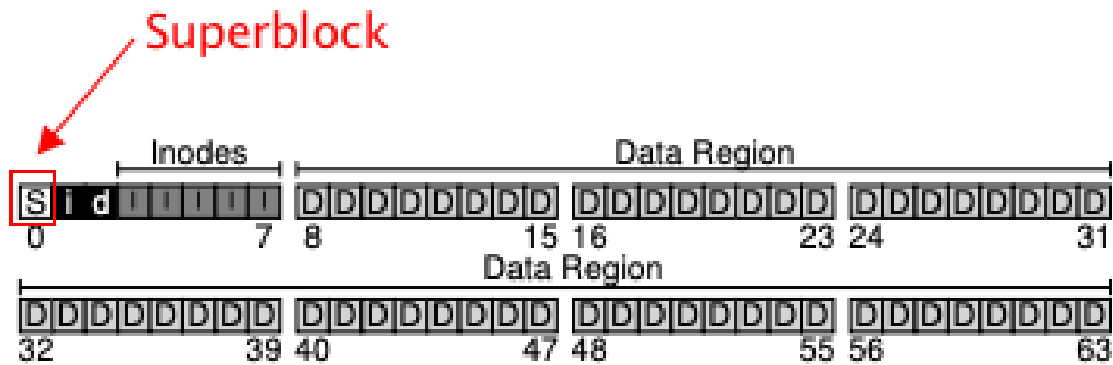


## 5 Index-based File System



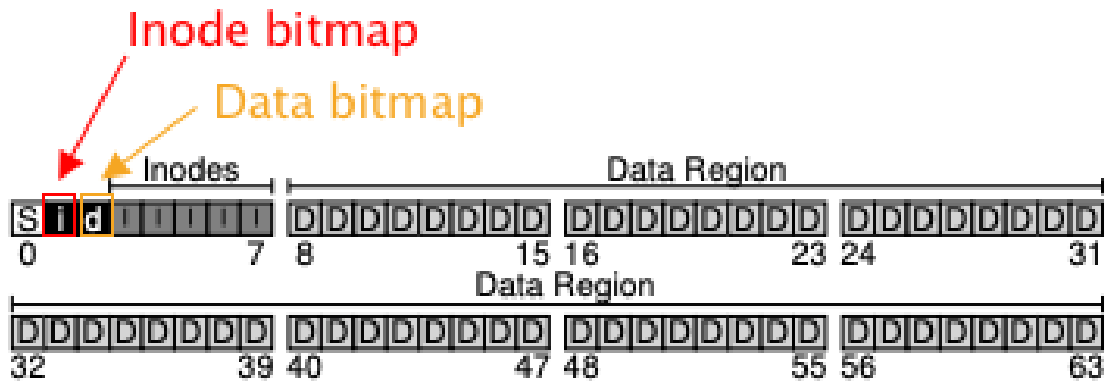
- Has following parts
  - Superblock
  - Inode Bitmap
  - Data Bitmap
  - Inodes
  - Data Region
- Each block in file system is 4KB
- Uses a large amount of metadata per file (especially for large files)

## 6 Superblock



- contains information about the file system, including
  - the number of inodes and data blocks in a particular file system
  - the magic number of some kind to identify the file system type (e.g NFS, FFS, VSFS)
- The OS reads superblock first to initialize various parameters, and then attach volume to the file-system tree

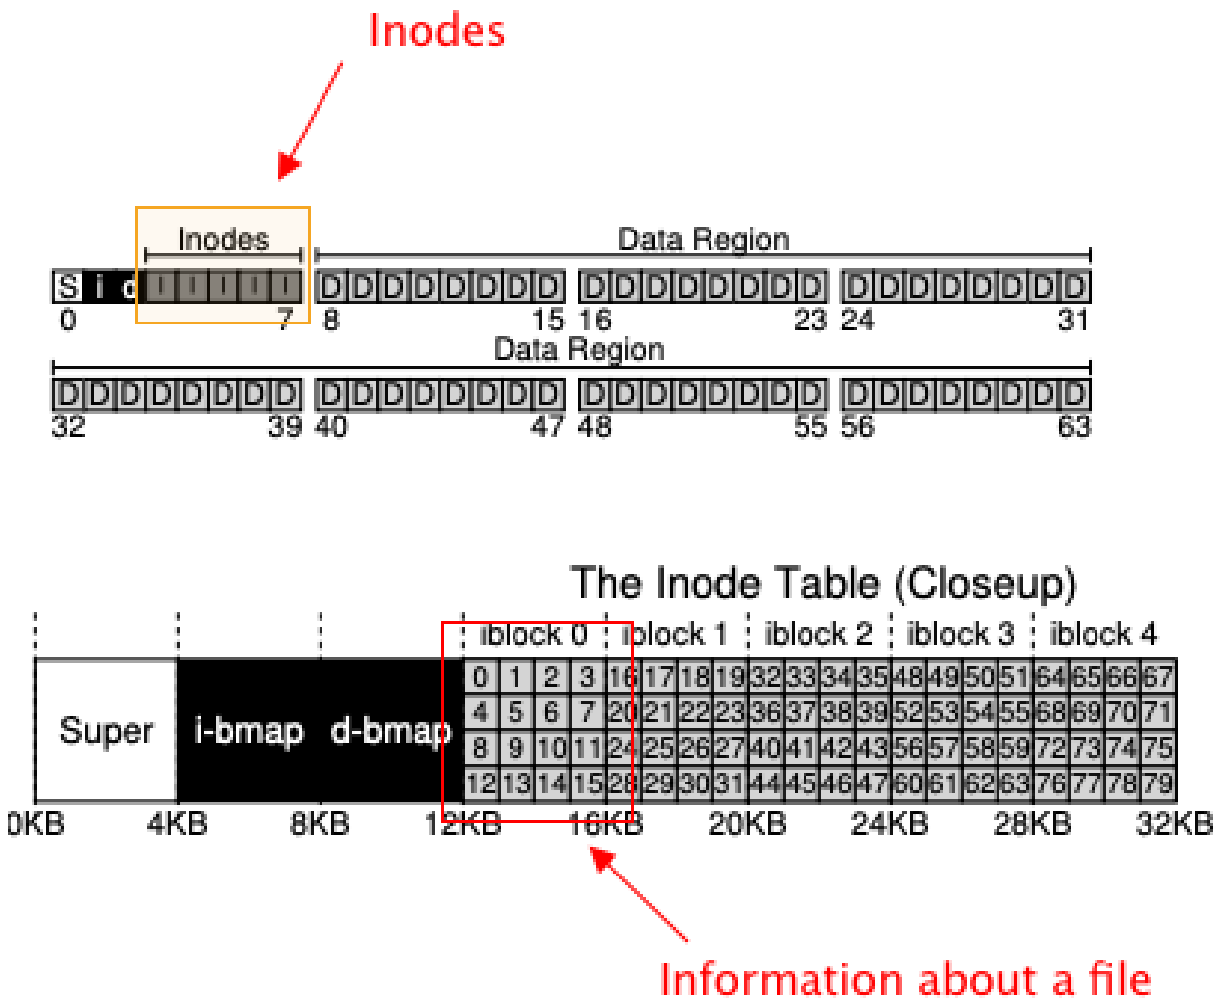
## 7 Bitmap



- Tracks whether inode or data blocks are free or allocated
- Is a simple and popular structure
- Uses each bit
  - 0 means free
  - 1 means in use

- **Data Bitmap** is bitmap for data region
- **Inode Bitmap** is bitmap for inode region

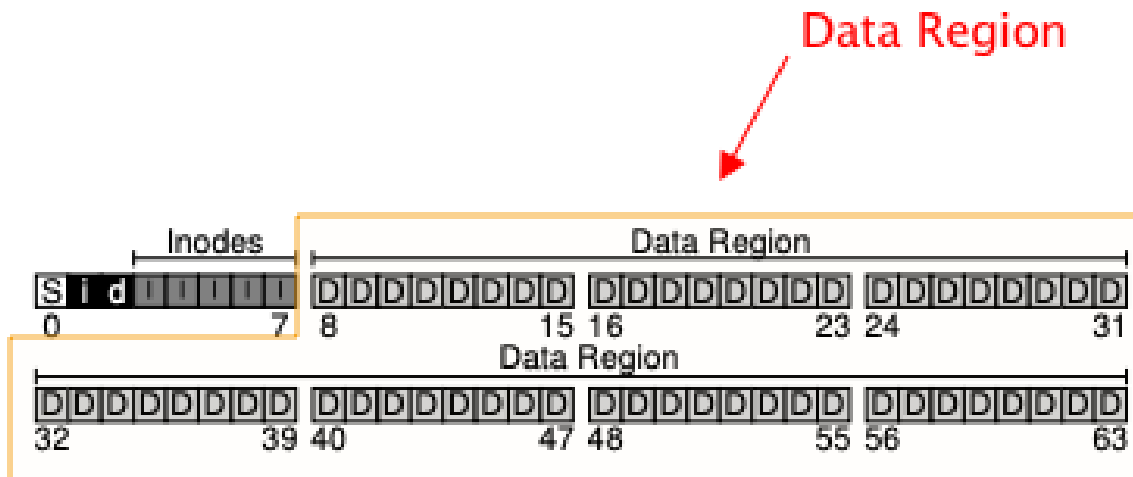
## 8 Inode



- Is a short form for **index node**
- Contains disk block location of the object's data <sup>[7]</sup>
- Contains all the information you need about a file (i.e. metadata)
  - File Type
    - \* e.g. regular file, directory, etc
  - Size
  - Number of blocks allocated to it
  - Protection information

- \* such as who owns the file, as well as who can access it
- Time information
  - \* e.g. When file was created, modified, or last accessed
- Location of data blocks reside on disk

## 9 Data Region



- Is the region of disk we use for user data

## 10 Block

- Size of each block is 4KB

## 11 lseek

- **Syntax:** `off_t lseek(int fildes, off_t offset, int whence)`
  - `fildes` - file descriptor
  - `offset` - file offset to a particular position in file

## 12 Kilobyte

- 1 kilobyte is 1024 bytes



## 13 file

- is an array of bytes which can be created, read, written and deleted
- low-level name is called **inode number** or **i-number**

## 14 Reading a File From Disk

### Example

When

```
open("/foo/bar", O_RDONLY)
```

is called

- the goal is to find the inode of the file `bar` to read its basic information (i.e. includes permission, information, file size etc)
- done by traversing the pathname and locate the desired inode
- Steps
  1. Begin traversal at the root of the file system, in the **root directory**
  2. Find **inode** of the root directory by looking for `i-number`
    - **i-number** is found in it's parent directoy
    - for root directory, there is no parent directory
    - it's inode number is 2 (for UNIX file systems)
  3. Read the **inode** of root directory
  4. Once its **inode** is read, look inside to find pointers to data blocks
  5. Recursively traverse the pathname until the desired inode is found (e.g `foo` → `bar`)
  6. Issue a `read()` system call to read from file
    - `fd` with offset 0 reads the first file block (e.g. `bar data[0]`)
    - `lseek(..., offset_amt * size_of_file_block)` is used to offset/move to desired block in `bar`
  7. Trasnfer data to `buf` data block
  8. Close `fd`. No I/O is read.

## 15 inode


- total size may vary
- inode pointer has size of 4 byte
- Has 12 **direct pointers** to 4KB data blocks
- Has 1 **indirect pointer** [when file grows large enough]
- Has 1 **double indirect Pointer** [when file grows large enough]
- Has 1 **triple indirect Pointer** [when file grows large enough]

## 16 Indirect Pointers

- Is allocated to data-block if file grows large enough
- Has total size of 4 KB or 4096 bytes
- Has  $4096/4 = 1024$  pointers
- Each pointer points to 4KB data-block
- File can grow to be  $(12 + 1024) \times 4K = 4144KB$

## 17 Double Indirect Pointers

- is allocated when single indirect pointer is not large enough
- each pointer in first pointer block points to another pointer block
- has  $1024^2$  pointers
- each of  $1024^2$  pointers point to 4KB data block
- File can grow to be  $(12 + 1024 + 1024^2) \times 4K = 4198448KB$  or  $\approx 4.20GB$



## 18 Triple Indirect Pointers

- is allocated when double indirect pointer is not large enough
- has  $1024^3$  pointers
- each of  $1024^3$  pointers point to 4KB data block
- File can grow to be  $(12 + 1024 + 1024^2 + 1024^3) \times 4K = 4299165744KB$  or  $\approx 4.00TB$

## 19 Old UNIX File system


- was simple, and looked like the following on disk




- has terrible performance
- suffers from **external fragmentation**
- had small data block (512 bytes) and transfer of data took too long

## 20 Fast File System

- modern file system has same APIS (`read()`, `write()`, `open()`, `close()`)
- divides into a number of **cylinder groups**



- each **block group** or **cylinder group** is consecutive portion of disk's address



## 21 Bitmap

- Are excellent way to manage free space
- tracks whether inodes/data block of the group are allocated

## 22 FFS Policies: Allocating Files and Directories


- Basic Idea: keep related stuff together, and keep related stuff far apart
- Directories Step
  - 1) Find the **cylinder group** with a low number of allocated directories and a high number of free inodes
    - low number of allocated directories → to balance directories across groups
    - high number of free nodes → to subsequently be able to allocate a bunch of files
  - 2) Put directory data and inode to the **cylinder group**
- Files Step
  - 1) Allocate the data blocks of a file in the same **cylinder group** as its inode
  - 2) Place all files in the same directory in the cylinder group of the directory they are in

### Example

On putting `/a/c`, `/a/d`, `/b/f`, FFS would place

- `/a/c`, `/a/d` as close as possible in the same **cylinder group**,
- `/b/f` located far away (in some other **cylinder group**)

## 23 Inode




- Is a short form for **index node**
- Contains disk block location of the object's data <sup>[1]</sup>
- Contains all the information you need about a file (i.e. metadata)
  - File Type
    - \* e.g. regular file, directory, etc
  - Size
  - Number of blocks allocated to it
  - Protection information
    - \* such as who owns the file, as well as who can access it
  - Time information
    - \* e.g. When file was created, modified, or last accessed
  - Location of data blocks reside on disk

## 24 Crash Consistency


- Goal: How to update persistent data structures despite the presence of a **power loss** or **system crash**?

## 25 Crash Scenarios

Before



After



1) Just the data block (Db) is written to disk

- No inode that points to it
- No bitmap that says the block is allocated
- It is as if the write never occurred
- There is no problem here. All is well. (In file system's point of view)

2) Just the updated inode (I[v2]) is written to disk

- Inode points to the disk where Db is about to be written
- No bitmap that says the block is allocated
- No Db is written
- Garbage data will be read
- Also creates **File-system Inconsistency**

- Caused by on-disk bitmap telling us Db 5 is not allocated, but inode saying it does

3) Just the updated bitmap ( $B[v2]$ ) is written to disk

- Bitmap indicates tht block 5 is allocated
- No inode exists at block 5
- Creates **file-system inconsistency**
- Creates **space-leak** if left as is
  - block 5 can never be used by the file system

4) Inode ( $I[v2]$ ) and bitmap ( $B[v2]$ ) are written to disk, and not data

- File system metadata is completely consistent (in perspective of file system)
- Garbage data will be read

5) Inode ( $I[v2]$ ) and data block (Db) are written, but not the bit map

- Creates **file-system inconsistency**
- Needs to be resolved before using file system again

6) Bitmap ( $B[v2]$ ) and data block (Db) are written, but not the inode ( $I[v2]$ )

- Creates **file-system inconsistency** between inode and data bitmap
- Creates **space-leak** if left as is
  - Inode block is lost for future use
- Creates **data-leak** if left as is
  - Data block is lost for future use

## 26 External Fragmentation


- Is various free holes that are generated in either your memory or disk space. <sup>[8]</sup>
- Are available for allocation, but may be too small to be of any use <sup>[8]</sup>



## 27 Internal Fragmentation

- Is wasted space within each allocated block <sup>[8]</sup>
- Occurs when more computer memory is allocated than is needed

## 28 Extent Based File System




- Is simply a disk pointer plus a length (in blocks)
  - Together, is called **extent**
- Often allows more than one extent
  - resolve problem of finding continuous free blocks
- Is less flexible but more compact
- Works well when there is enough free space on the disk and files can be laid out contiguously

### Example

Linux's ext4 file system

## 29 Fields

- Is the members in a structure



## 30 Process List

- Is a data structure in kernel or OS
- Contains information about all the processes running in the system

## 31 Process Control Block

- Is a data structure in kernel or OS
- Contains all information about a process
- Is where the OS keeps all of a process' hardware execution state
- Generally includes
  1. Process state (ready, running, blocked)
  2. Process number
  3. Program counter: address of the next instruction
  4. CPU Registers: is saved at an interrupt
  5. CPU scheduling information: process priority
  6. Memory management info: page tables
  7. I/O status information: list of open files